

SHRI NARAYAN PANDEY

Software Engineer | Backend | Microservices | Generative AI

shrinarayanp03@gmail.com | +91 9045893362 | Hyderabad, India | github.com/ShriNarayanPandey01 | linkedin.com/in/shrinarayanpandey01

PROFESSIONAL SUMMARY

Software Engineer with hands-on experience in backend development, microservice architecture, and Generative AI. Currently at HashedIn by Deloitte, designing and building enterprise-grade systems using FastAPI and Google ADK. Contributed schema design across 4 microservices in a 10+ microservice platform and implemented Redis-based session management to reduce database load and improve latency. Skilled in Agile/Scrum workflows, LLD, and peer code review practices. Built an AI-powered research assistant during internship using LangChain, LangGraph, FastAPI, and GCP.

CORE COMPETENCIES

FastAPI | Python | Java | Microservices | REST APIs | LangChain | LangGraph | Google ADK | RAG | LLM Pipelines | Redis | PostgreSQL | Apache Kafka | Docker | GCP | Spring Boot | Agile/Scrum | Low-Level Design | System Design | DSA

TECHNICAL SKILLS

Languages: Python, Java, JavaScript, Bash

Backend Frameworks: FastAPI, Flask, Spring Boot, Node.js

GenAI & Agentic AI: Google ADK, LangChain, LangGraph, Prompt Engineering, RAG, LLM Pipelines

Databases & Caching: PostgreSQL, MySQL, Redis (session management), SQLAlchemy

Microservices: REST APIs, Schema Design, Apache Kafka, Inter-service Communication

DevOps & Infrastructure: Docker, Linux, GCP, Git, Bash Scripting, CI/CD

Methodologies: Agile, Scrum, Low-Level Design (LLD), OOP, System Design, Data Structures & Algorithms

Frontend: Angular, React.js

WORK EXPERIENCE

Software Engineer 1 | [HashedIn by Deloitte](#) | Hyderabad, India | Dec 2025 – Present

- Developed and maintained backend services using FastAPI and Google ADK for a confidential enterprise platform, covering API design, agentic workflow integration, and core business logic.
- Designed schema across 4 microservices in a 10+ microservice system — defined entity models, table relationships, API contracts, and service boundaries, reducing cross-service coupling and maintaining clear data ownership.
- Implemented Redis-based session management across microservices, eliminating repeated database calls for authenticated users and measurably improving session lookup latency.
- Participated in fortnightly Agile/Scrum sprints; contributed to LLD discussions and peer code reviews to uphold code quality and engineering standards.

Software Developer Intern | [HashedIn by Deloitte](#) | Bengaluru, India | Oct 2025 – Dec 2025

- Completed structured training across FastAPI, Spring Boot, Java, LangChain, LangGraph, Angular, React, and Node.js, with graded hands-on assignments in each area.
- Built REST APIs and backend applications using Spring Boot and FastAPI as training deliverables; deployed and ran services using Docker on Linux environments.
- Led backend development of an AI-powered research assistant as internship capstone — collaborated with 3 fellow interns, architecting backend modules using FastAPI, LangChain, and LangGraph.
- Attended code reviews and LLD walkthroughs led by senior engineers, gaining hands-on experience with real-world Agile engineering practices.
- Converted to full-time Software Engineer 1 on successful completion of the training program.

Software Engineer Intern | [Telaverge Communications](#) | Bengaluru, India | Dec 2024 – Jun 2025

- Assisted with backend development and API testing for telecom software components running on Linux-based infrastructure.

- Automated 10+ repetitive operational tasks using Python and Bash scripts, reducing manual effort and accelerating issue resolution cycles.
- Gained experience with enterprise SDLC workflows, production support processes, and distributed system troubleshooting alongside senior engineers.

PROJECTS

AI Research Assistant — *FastAPI · LangChain · LangGraph · Google ADK · GCP · Angular* | Team of 4

- Built an AI-powered platform that fetches and summarizes the latest news on user-defined topics to support in-depth research using Generative AI and agentic orchestration.
- Architected the agentic workflow, service boundaries, and data flow in early design phases; developed the FastAPI backend and LangChain/LangGraph orchestration modules end-to-end.
- Integrated Google ADK for external tool-calling; coordinated GCP deployment and Angular frontend integration, delivering a working product within the internship timeline.

Kafka from Scratch — *Java · Distributed Systems · Binary Protocol · Concurrency* [GitHub](#)

- Re-implemented core Kafka internals in Java — binary wire-protocol parsing, producer/consumer request-response cycles, topic-metadata flow, and log-segment management from first principles.
- Applied OS-level concepts including multi-threading, socket programming, and I/O buffering to sustain reliable throughput across concurrent producer and consumer streams.
- Built practical understanding of partition leadership, offset tracking, and consumer-group coordination — directly applicable to scalable distributed backend design.

Custom Language Interpreter — *Python · Compiler Design · AST · Runtime Engine* [GitHub](#)

- Engineered a complete interpreter pipeline in Python: hand-written lexer for tokenization, recursive-descent parser producing an AST, and a tree-walking runtime for execution.
- Runtime supports variable scoping with environment chaining, control flow (if/else, while loops), and user-defined functions with a proper call stack for non-trivial programs.

AI Travel Assistant — *Python · LangChain · Conversational Memory · Tool Calling* [GitHub](#)

- Built an LLM-powered assistant for hotel/flight suggestions and itinerary planning using LangChain agent workflows with conversational memory to retain context across turns.
- Designed prompt pipeline and tool-calling layer so the agent retrieves live data rather than relying on model knowledge, producing grounded, multi-turn recommendations.

EDUCATION

B.Tech — Computer Science (AI & ML Specialization) | Graduated 2025

Pranveer Singh Institute of Technology, Kanpur, India

ACHIEVEMENTS

- Ranked 88th out of 50,000+ teams in the Amazon ML Challenge 2024.
- Achieved 1700+ rating on LeetCode, consistently placing in the top percentile across weekly and biweekly contests.
- Converted from Software Developer Intern to full-time Software Engineer 1 at HashedIn by Deloitte on successful completion of the training program.